

Machine Learning III

Dr. Angelica Lim

Assistant Professor

School of Computing Science

Simon Fraser University, Canada

Oct. 1, 2024

Course Overview

Week 1 : Getting to know you

Week 2 : Introduction to Artificial Intelligence

Week 3 : Machine Learning I: Basic Supervised Models (Classification)

Week 4 : Machine Learning II: Supervised Regression, Classification and Gradient Descent, K-Means

Week 5 : Machine Learning III: Neural Networks and Backpropagation

Week 6 : Search (A*)

Week 7 : Markov Decision Processes

Week 8 : Midterm

Week 9 : Reinforcement Learning

Week 10 : Games

Week 11 : Hidden Markov Models and Bayesian Networks

Week 12 : Constraint Satisfaction Problems

Week 13 : Ethics and Explainability



Today's Plan

Supervised Learning

19.6.2 - Stochastic Gradient Descent

19.6.3 - Multivariable Linear Regression

19.4.3 - Dimensionality and Regularization

21.1 - Neural Networks

21.1.2 - Backpropagation

Course Announcements

Assignment 1 is released and due October 7

No class on Tues, October 15 due to holiday as per SFU guidelines


Midterm is set for Tues, October 29 in class.

Supervised Learning

Stochastic Gradient Descent

Gradient Descent is Slow

$$\text{TrainLoss}(\mathbf{w}) = \frac{1}{|\mathcal{D}_{\text{train}}|} \sum_{(x,y) \in \mathcal{D}_{\text{train}}} \text{Loss}(x, y, \mathbf{w})$$



Algorithm: gradient descent

Initialize $\mathbf{w} = [0, \dots, 0]$

For $t = 1, \dots, T$:

$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} \text{TrainLoss}(\mathbf{w})$

We only update \mathbf{w} after finding the average over all samples

Stochastic gradient descent



Algorithm: stochastic gradient descent

Initialize $\mathbf{w} = [0, \dots, 0]$

For $t = 1, \dots, T$:

For $(x, y) \in \mathcal{D}_{\text{train}}$:

$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} \text{Loss}(x, y, \mathbf{w})$

We update \mathbf{w} after finding the loss of each sample

Variants of SGD include 1) minibatch, to update using an average over B samples, 2) randomizing the order over the samples. (*Why would this be important? What if you had all the positive samples first, followed by negative samples?*)

For assignment, you have to think if the feature will be useful considering the data you have, and whether it will scale well.

Supervised Learning

Nonlinear Functions

Recall: Linear Regression

training data

x	y
1	1
2	3
4	3

learning algorithm



2.71

Which predictors are possible?

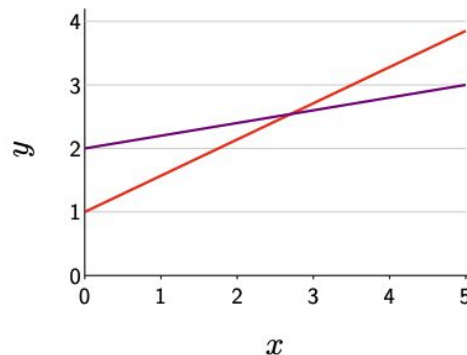
Hypothesis class

$$\mathcal{F} = \{f_{\mathbf{w}}(x) = \mathbf{w} \cdot \phi(x) : \mathbf{w} \in \mathbb{R}^d\}$$


$$\phi(x) = [1, x]$$

$$f(x) = [1, 0.57] \cdot \phi(x)$$

$$f(x) = [2, 0.2] \cdot \phi(x)$$



More complex data



How do we fit a non-linear predictor?

From linear to quadratic predictors

We can get a non-linear predictor just by changing the feature extractor $\rightarrow \phi(x) = [1, x, x^2]$ e.g. $\phi(3) = [1, 3, 9]$


$$f(x) = [2, 1, -0.2] \cdot \phi(x)$$

$$f(x) = [4, -1, 0.1] \cdot \phi(x)$$

$$f(x) = [1, 1, 0] \cdot \phi(x)$$

If we set x^2 to zero, we recover a linear predictor

$$\mathcal{F} = \{f_{\mathbf{w}}(x) = \mathbf{w} \cdot \phi(x) : \mathbf{w} \in \mathbb{R}^3\}$$



Piecewise constant predictors


We can also obtain a non-linear predictor by partitioning the input space

$$\phi(x) = [\mathbf{1}[0 < x \leq 1], \mathbf{1}[1 < x \leq 2], \mathbf{1}[2 < x \leq 3], \mathbf{1}[3 < x \leq 4], \mathbf{1}[4 < x \leq 5]]$$

$$f(x) = [1, 2, 4, 4, 3] \cdot \phi(x)$$

$$f(x) = [4, 3, 3, 2, 1.5] \cdot \phi(x)$$

$$\mathcal{F} = \{f_{\mathbf{w}}(x) = \mathbf{w} \cdot \phi(x) : \mathbf{w} \in \mathbb{R}^5\}$$




Predictors with periodicity

$$\phi(x) = [1, x, x^2, \cos(3x)] \quad \text{e.g. } \phi(2) = [1, 2, 4, 0.96]$$

$$f(x) = [1, 1, -0.1, 1] \cdot \phi(x)$$


$$f(x) = [3, -1, 0.1, 0.5] \cdot \phi(x)$$

$$\mathcal{F} = \{f_{\mathbf{w}}(x) = \mathbf{w} \cdot \phi(x) : \mathbf{w} \in \mathbb{R}^4\}$$



What prediction task might be periodic?

Many hypothesis classes are possible!



Linear in what?

$$f_{\mathbf{w}}(x) = \mathbf{w} \cdot \phi(x)$$

Linear in \mathbf{w} ? Yes

Linear in $\phi(x)$? Yes

Linear in x ? No!

Indeed, the raw values of our data may not even be numbers

$\mathbf{w} \cdot \phi(x)$ can be a **non-linear** function of x


$\mathbf{w} \cdot \phi(x)$ always a **linear** function of \mathbf{w}

If the **function output** $\mathbf{w} \cdot \phi(x)$ is **linear in \mathbf{w}** and the **loss function is convex** (squared, hinge, logistic losses but not zero-one loss), then minimizing the training loss with gradient descent and a proper step size is guaranteed to converge to the global minimum.

Recall: Linear classification

$$\phi(x) = [x_1, x_2]$$

$$f(x) = \text{sign}([-0.6, 0.6] \cdot \phi(x))$$



The decision boundary is a line

From linear to quadratic classifiers

Add a new feature $x_1^2 + x_2^2$
into the feature vector

$$\phi(x) = [x_1, x_2, x_1^2 + x_2^2]$$


$$f(x) = \text{sign}([2, 2, -1] \cdot \phi(x))$$

Equivalently:

$$f(x) = \begin{cases} 1 & \text{if } \{(x_1 - 1)^2 + (x_2 - 1)^2 \leq 2\} \\ -1 & \text{otherwise} \end{cases}$$

What is the result
of $x = [0, 0]$?

The decision boundary is a circle



Visualization in feature space


Input space: $x = [x_1, x_2]$, decision boundary is a circle

Feature space: $\phi(x) = [x_1, x_2, x_1^2 + x_2^2]$, decision boundary is a hyperplane

Want to create a separator where the points inside a circle belong to one group and the points outside belong to another group.

If your initial feature space is not expressive enough, projecting into a higher dimension will allow you to linearly separate your data points.

Decision trees are very expressive.
For decision trees, you need to handcraft many of your features.




Summary

$$f_{\mathbf{w}}(x) = \mathbf{w} \cdot \phi(x)$$

linear in \mathbf{w} , $\phi(x)$

non-linear in x



Regression: non-linear predictor, classification: non-linear decision boundary

Types of non-linear features: quadratic, piecewise constant, etc.

Supervised Learning

Neural Networks

Motivating example




Example: predicting car collision

Input: positions of two oncoming cars $x = [x_1, x_2]$

Output: whether safe ($y = +1$) or collide ($y = -1$)

x_1	x_2	y
0	2	1
2	0	1
0	0	-1
2	2	-1



Safe if cars are sufficiently far:

$$y = \text{sign}(|x_1 - x_2| - 1)$$

The famous XOR problem that is impossible to fit using a linear classifier

Decomposing the problem

Test if car 1 is far right of car 2:

$$h_1(x) = 1[x_1 - x_2 \geq 1]$$


Test if car 2 is far right of car 1:

$$h_2(x) = 1[x_2 - x_1 \geq 1]$$

Safe if at least one is true:

$$f(x) = \text{sign}(h_1(x) + h_2(x))$$

x	$h_1(x)$	$h_2(x)$	$f(x)$
$[0, 2]$	0	1	+1
$[2, 0]$	1	0	+1
$[0, 0]$	0	0	-1
$[2, 2]$	0	0	-1



Rewriting using vector notation

$$h_1(x) = \mathbf{1}[x_1 - x_2 \geq 1]:$$

$$h_1(x) = \mathbf{1}[x_1 - x_2 \geq 1]$$

$$\rightarrow \mathbf{1}[-1 + x_1 - x_2 \geq 0]$$

$$\rightarrow \mathbf{1}[\mathbf{[-1, 1, -1]} \cdot [1, x_1, x_2] \geq 0]$$

$$\mathbf{h}(x) = \mathbf{1} \left[\begin{bmatrix} -1 & +1 & -1 \\ -1 & -1 & +1 \end{bmatrix} \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix} \geq 0 \right]$$

v1

v2

$$h_2(x) = \mathbf{1}[x_2 - x_1 \geq 1]$$

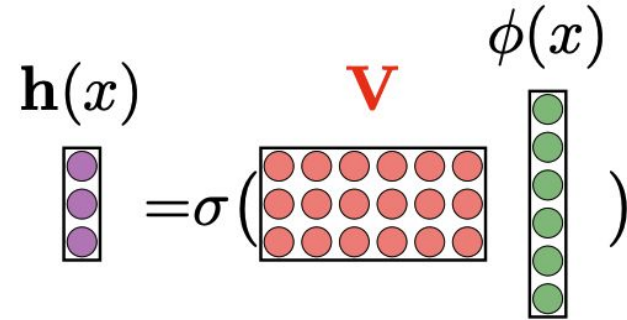
$$h_2(x) = \mathbf{1}[x_2 - x_1 \geq 1]$$

$$\rightarrow \mathbf{1}[-1 - x_1 + x_2 \geq 0]$$

$$\rightarrow \mathbf{1}[\mathbf{[-1, -1, 1]} \cdot [1, x_1, x_2] \geq 0]$$

$$f(x) = \text{sign}(h_1(x) + h_2(x)) = \text{sign}(\mathbf{[1, 1]} \cdot \mathbf{h}(x))$$

Two-layer neural networks


$$\mathbf{h}(x) = \sigma(\mathbf{V} \phi(x))$$


Avoiding zero gradients

Problem: gradient of $h_1(x)$ with respect to \mathbf{v}_1 is 0

$$h_1(x) = \mathbf{1}[\mathbf{v}_1 \cdot \phi(x) \geq 0]$$

Solution: replace with an **activation function** σ with non-zero gradients



- Threshold: $\mathbf{1}[z \geq 0]$
- Logistic: $\frac{1}{1+e^{-z}}$
- ReLU: $\max(z, 0)$

$$h_1(x) = \sigma(\mathbf{v}_1 \cdot \phi(x))$$

Two-layer neural networks

Intermediate subproblems:

$$\mathbf{h}(x) = \sigma \left(\mathbf{V} \phi(x) \right)$$

Predictor (classification):

$$f_{\mathbf{V}, \mathbf{w}}(x) = \text{sign} \left(\mathbf{w} \cdot \mathbf{h}(x) \right)$$

We can interpret $\mathbf{h}(x)$ as a learned feature representation!

Deep neural networks

1-layer network:

$$\text{score} = \mathbf{w} \cdot \phi(x)$$

2-layer network:

$$\text{score} = \mathbf{w} \cdot \sigma \left(\mathbf{V} \cdot \phi(x) \right)$$


3-layer network:

$$\text{score} = \mathbf{w} \cdot \sigma \left(\mathbf{V}_2 \cdot \sigma \left(\mathbf{V}_1 \cdot \phi(x) \right) \right)$$


Layers represent multiple layers of abstraction



3rd layer
"Objects"



2nd layer
"Object parts"



1st layer
"Edges"




Pixels

Summary


Linear predictors:

$$f_{\mathbf{w}}(x) = \mathbf{w} \cdot \phi(x), \quad \phi(x) = [1, x]$$



Non-linear (quadratic) predictors:


$$f_{\mathbf{w}}(x) = \mathbf{w} \cdot \phi(x), \quad \phi(x) = [1, x, x^2]$$



V is like w, but it allows for multiple subproblems.

Non-linear neural networks:


$$f_{\mathbf{w}}(x) = \mathbf{w} \cdot \sigma(\mathbf{V}\phi(x)), \quad \phi(x) = [1, x]$$



Artificial neural networks



They are inspired from actual neurons in human/animal brain.

- **Dendrites:** Take its input from other neurons in the form of electrical impulses.
- **Soma:** Generates inferences from inputs and decides what action to take.
- **Axon terminals:** Transmit outputs in the form of impulses.




Artificial neuron cell (aka perceptron)

$$a_j = \sum_{i=1}^D (w_{ji}^{(1)} x_i + x_{j0}^{(1)})$$



Artificial neuron cell (aka perceptron)



Backpropagation

Motivation: regression with 4-layer neural network

Loss on one example:

$$\text{Loss}(x, y, \mathbf{V}_1, \mathbf{V}_2, \mathbf{V}_3, \mathbf{w}) = (\mathbf{w} \cdot \sigma(\mathbf{V}_3 \sigma(\mathbf{V}_2 \sigma(\mathbf{V}_1 \phi(x)))) - y)^2$$

(Stochastic) gradient descent:

$$\mathbf{V}_1 \leftarrow \mathbf{V}_1 - \eta \nabla_{\mathbf{V}_1} \text{Loss}(x, y, \mathbf{V}_1, \mathbf{V}_2, \mathbf{V}_3, \mathbf{w})$$

$$\mathbf{V}_2 \leftarrow \mathbf{V}_2 - \eta \nabla_{\mathbf{V}_2} \text{Loss}(x, y, \mathbf{V}_1, \mathbf{V}_2, \mathbf{V}_3, \mathbf{w})$$

$$\mathbf{V}_3 \leftarrow \mathbf{V}_3 - \eta \nabla_{\mathbf{V}_3} \text{Loss}(x, y, \mathbf{V}_1, \mathbf{V}_2, \mathbf{V}_3, \mathbf{w})$$

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} \text{Loss}(x, y, \mathbf{V}_1, \mathbf{V}_2, \mathbf{V}_3, \mathbf{w})$$

This can get tedious!

Computation graphs

$$\text{Loss}(x, y, \mathbf{V}_1, \mathbf{V}_2, \mathbf{V}_3, \mathbf{w}) = (\mathbf{w} \cdot \sigma(\mathbf{V}_3 \sigma(\mathbf{V}_2 \sigma(\mathbf{V}_1 \phi(x)))) - y)^2$$



Definition: computation graph


A directed acyclic graph whose root node represents the final mathematical expression and each node represents intermediate subexpressions.

Upshot: compute gradients via general **backpropagation** algorithm


Purposes:

- Automatically compute gradients (how TensorFlow and PyTorch work)
- Gain insight into modular structure of gradient computations

Functions as boxes



Basic building blocks




Linear classification with hinge loss


$$\text{Loss}(x, y, \mathbf{w}) = \max\{1 - \mathbf{w} \cdot \phi(x)y, 0\}$$

$$\nabla_{\mathbf{w}} \text{Loss}(x, y, \mathbf{w}) = -1[\text{margin} < 1]\phi(x)y$$

Can use the computation graph to obtain the gradient



Text



Just the product of the gradients on the path from \mathbf{w} to the root!

Two-layer neural networks

Recall squared loss:



$$\text{Loss}(x, y, \mathbf{w}) = (f_{\mathbf{w}}(x) - y)^2$$

$$\text{Loss}(x, y, \mathbf{V}, \mathbf{w}) = (\mathbf{w} \cdot \sigma(\mathbf{V}\phi(x)) - y)^2$$

$$\nabla_{\mathbf{w}} \text{Loss}(x, y, \mathbf{V}, \mathbf{w}) = 2(\text{residual})\mathbf{h}$$

$$\nabla_{\mathbf{V}} \text{Loss}(x, y, \mathbf{V}, \mathbf{w}) = 2(\text{residual})\mathbf{w} \circ \mathbf{h} \circ (1 - \mathbf{h})\phi(x)^{\top}$$

Review computation graph



\circ is element-wise multiplication

Backpropagation

$$\text{Loss}(x, y, \mathbf{w}) = (\mathbf{w} \cdot \phi(x) - y)^2$$

$$\mathbf{w} = [3, 1], \phi(x) = [1, 2], y = 2$$



Algorithm: backpropagation algorithm

Forward pass: compute each f_i (from leaves to root)

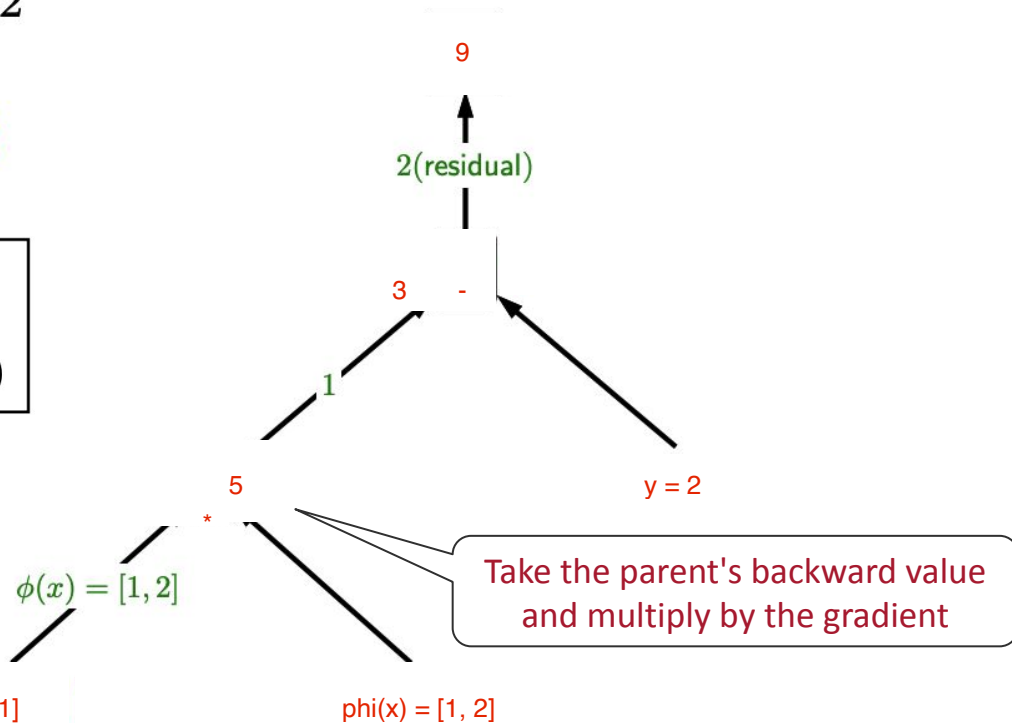
Backward pass: compute each g_i (from root to leaves)



Definition: Forward/backward values


Forward: f_i is value for subexpression rooted at i

Backward: $g_i = \frac{\partial \text{loss}}{\partial f_i}$ is how f_i influences loss



Machine Learning Strategies

How do you prevent overfitting?



How good is the predictor f ?



Key idea: the real learning objective

Our goal is to minimize **error on unseen future examples**.

As the hypothesis class size increases, the estimation error increases

Strategy 1: Reduce Dimensionality

Manual feature selection:

- Add features if they help
- Remove features if they don't help

Automatic feature selection (beyond the scope of this class):

- Boosting

Strategy 2: Regularization

L2 regularization penalizes the norm (length) of \mathbf{w} by λ .

$$\min_{\mathbf{w}} \text{TrainLoss}(\mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$



Algorithm: gradient descent

Initialize $\mathbf{w} = [0, \dots, 0]$

For $t = 1, \dots, T$:

$$\mathbf{w} \leftarrow \mathbf{w} - \eta(\nabla_{\mathbf{w}} \text{TrainLoss}(\mathbf{w}) + \lambda \mathbf{w})$$

Note: Support vector machines are exactly hinge loss + L2 regularization